Nonparametric-efficient Causal Mediation Analysis for Stochastic Interventions

Nima Hejazi, Mark van der Laan, and Iván Díaz

Graduate Group in Biostatistics & Dept. of Statistics, UC Berkeley

Division of Biostatistics, Dept. of Healthcare Policy & Research, Weill Cornell Medicine

Overview & Motivations

- Using stochastic interventions, we present a decomposition of the *population intervention effect* into direct and indirect effects.
  - Define causal contrasts of effects of continuous and categorical exposures.
  - Identify a parameter required to construct both direct and indirect effects.
- We propose estimators for constructing these direct and indirect effects:
  - *Classical*: substitution and IPW estimators using parametric methods.
  - *Nonparamtric-efficient*: one-step and TMLE leveraging machine learning.
- Our efficient estimators are asymptotically linear under $n^{1/4}$-consistency of nuisance functions (may use highly adaptive lasso).

Software Implementation

- The `medshift` R package [3] implements these estimators and leverages state-of-the-art machine learning in the procedure.
  - Construction of all estimators via the eponymous `medshift()` function.
  - Uses the `sl3` R package to incorporate machine learning facilities.
  - Relies on the framework of the `tmle3` R package for TMLE implementation.
  - Flexible cross-fitting implementation via the `origami` R package.
- `sl3`, `tmle3`, and `origami` are core engines of the new **`tlverse`** software ecosystem.
  - Check out `https://tlverse.org`
  - Our handbook: `https://tlverse.org/tlverse-handbook`

Stochastic Population Intervention (In)Direct Effects

- Consider $O = (W, A, Z, Y) \sim P_0 \in \mathcal{M}$, for W a set of baseline covariates, A an intervention, Y the outcome, and Z a mediator between A and Y, with no assumptions on nonparametric model $\mathcal{M}$.
- We decompose the total population intervention effect (PIE) in terms of a *population intervention direct effect (PIDE)* and a *population intervention indirect effect (PIIE)*:

$$\psi(\delta) = \overbrace{\mathbb{E}\{Y(g, q) - Y(g_\delta, q)\}}^{\text{PIDE}} + \overbrace{\mathbb{E}\{Y(g_\delta, q) - Y(g_\delta, q_\delta)\}}^{\text{PIIE}}.$$

- We show the causal parameter $\mathbb{E}\{Y(g_\delta, q)\}$ is identified by the observed data parameter [2]:

$$\theta(\delta) = \int m(z, a, w) g_\delta(a \mid w) p(z, w) d\nu(a, z, w),$$

for outcome mechanism $m(z, a, w)$ and post-intervention treatment mechanism $g_\delta(a \mid w)$ — a stochastic intervention on A while drawing Z as though A came from its natural distribution.

- Letting $\eta = (g, e, m, \phi)$, the efficient influence function for $\theta(\delta)$ in the nonparametric model $\mathcal{M}$ is $D^Y_{\eta,\delta}(o) + D^A_{\eta,\delta}(o) + D^{Z,W}_{\eta,\delta}(o) - \theta(\delta)$ for

$$D^{Z,W}_{\eta,\delta}(o) = \int m(z, a, w) g_\delta(a \mid w) d\kappa(a), \quad D^Y_{\eta,\delta}(o) = \frac{\mathbf{g_\delta(a \mid w)}}{\mathbf{e(a \mid z, w)}}\{y - m(z, a, w)\},$$

$$D^A_{\eta,\delta}(o) = \frac{\delta\phi(\mathbf{w})\{a - g(1 \mid w)\}}{\{\boldsymbol{\delta}\mathbf{g(1 \mid w) + 1 - g(1 \mid w)}\}^{\mathbf{2}}}, \quad \phi(w) = \mathbb{E}\{m(1, Z, W) - m(0, Z, W) \mid W = w\},$$

where, for simplicity, we present the case $A \in \{0, 1\}$. For an unabridged treatment, see [2].

Construction of Nonparametric-Efficient Estimators

- To avoid entropy conditions on initial estimators, we rely on cross-validation [6, 1], letting $\hat{\eta}_j$ be the estimator of $\eta = (g, e, m, \phi)$ and $j(i)$ the index of the validation set containing observation i.
- A one-step estimator [4] may be constructed by augmenting the substitution estimator with the auxiliary scores (D^A and D^Y) in the efficient influence function (EIF), yielding

$$\hat{\theta}_{\text{OS}}(\delta) = \frac{1}{n}\sum_{i=1}^{n} D_{\hat{\eta}_{j(i)},\delta}(O_i) = \frac{1}{n}\sum_{i=1}^{n}\left\{D^Y_{\hat{\eta}_{j(i)},\delta}(O_i) + D^A_{\hat{\eta}_{j(i)},\delta}(O_i) + D^{Z,W}_{\hat{\eta}_{j(i)},\delta}(O_i)\right\}.$$

- A targeted minimum loss estimator (TMLE) may be constructed by using the efficient influence function to update components of the substitution estimator via a targeting procedure:

$$\hat{\theta}_{\text{TMLE}}(\delta) = \int m^\star_n(z, a, w) g^\star_{\delta,n}(a \mid w) d\kappa(a),$$

where $g^\star_{\delta,n}(a \mid w)$ and $m^\star_n(z, a, w)$ come from *targeting* processes using parametric models $\text{logit}(g_{\delta,n,k\epsilon}) = g_{\delta,n,(k-1)\epsilon} + \epsilon_\Delta \mathbf{H}^{\mathbf{A}}_{\mathbf{(k-1)}\epsilon}$ and $\text{logit}(m_{n,k\epsilon}) = m_{n,(k-1)\epsilon} + \epsilon_\Delta \mathbf{H}^{\mathbf{Y}}_{\mathbf{(k-1)}\epsilon}$, respectively.

  - Unlike the one-step estimator, TMLE constructs a substitution estimator, respecting bounds.
  - Targeting step uses the method of unverisal least favorable one-dimensional submodels [5].
- Both are multiply robust and allow confidence intervals and hypothesis tests to be constructed from an EIF-based variance estimator $\mathbb{V}(\hat{\theta}(\delta)) = D^2_{\hat{\eta},\delta}(O)$, valid even when using machine learning.

Results & Discussion

- All estimators approximately unbiased in large samples; TMLE outperforms one-step uniformly in smaller-sample settings.
- Inference is only valid for the one-step and TMLE when using machine learning (here, HAL) for estimating nuisance regressions.

References

[1] V. Chernozhukov, D. Chetverikov, M. Demirer, E. Duflo, C. Hansen *et al.*, "Double machine learning for treatment and causal parameters," *arXiv preprint arXiv:1608.00060*, 2016.

[2] I. Díaz and N. S. Hejazi, "Causal mediation analysis for stochastic interventions," *in revision*, 2019. [Online]. Available: https://arxiv.org/abs/1901.02776

[3] N. S. Hejazi and I. Díaz, *medshift: Causal mediation analysis for stochastic interventions in R*, 2019, R package version 0.0.8. [Online]. Available: https://github.com/nhejazi/medshift

[4] J. Pfanzagl and W. Wefelmeyer, "Contributions to a general asymptotic statistical theory," *Statistics & Risk Modeling*, vol. 3, no. 3-4, pp. 379–388, 1985.

[5] M. van der Laan and S. Gruber, "One-step targeted minimum loss-based estimation based on universal least favorable one-dimensional submodels," *The international journal of biostatistics*, vol. 12, no. 1, pp. 351–378, 2016.

[6] W. Zheng and M. J. van der Laan, "Cross-validated targeted minimum-loss-based estimation," in *Targeted Learning*. Springer, 2011, pp. 459–474.

But wait, there's more!

- **N. Hejazi**: `nhejazi@berkeley.edu`; **M. van der Laan**: `laan@berkeley.edu`; **I. Díaz**: `ild2005@med.cornell.edu`
- `https://arxiv.org/abs/1901.02776`
- Check out Iván's talk Friday morning!